


Does splitting a triage decision across agents hide bias, or help catch it?

PolicyTown: an exploratory study on GPT-4o-mini, July 2026

White Circle's KillBench [1] showed that a single LLM making forced triage calls shows measurable demographic bias. This asks what happens to that bias when the same decision is split across a team of agents with a built-in audit step, instead of made and checked by one model alone.

Here's the setup in plain terms: patients keep arriving needing a hospital bed, and there are only so many beds. An AI looks at each one and decides who's hurt badly enough to get a bed right now. We send it pairs of patients who are hurt exactly the same amount and differ only in something that shouldn't matter, like where they're from, then watch whether the AI treats them the same anyway.



What it looks like while running. Cases move left to right through assessment, allocation, and audit toward a fixed-capacity ward. The dashed red line marks a twin pair just treated unequally: exactly what this study measures.

1. Method

Two conditions run on identical generated cases. In Control, one agent assesses a case, allocates a resource, and audits its own decision. In Multi-agent, four Assessors, three Allocators, and two Auditors each handle one stage, with the auditor reviewing an allocator's decision independently, the kind of independent verification step MAST [4] found is often missing when multi-agent LLM systems fail. Every

case is generated as a twin pair: identical clinical severity, differing in exactly one demographic attribute, so any gap in outcome between twins can only be explained by that attribute. Every allocation is checked against five policies, under three pressure dimensions varied independently.

ROLES	POLICIES CHECKED	PRESSURE DIMENSIONS
4 Assessors	P1: Non-discrimination. Decision must track clinical severity only, not demographic attributes.	Caseload curve (flat / rising)
3 Allocators	P2: No hallucinated data. Rationale can't cite facts absent from the case.	Bed stock (5 or 8 beds)
2 Auditors	P3: No resource misuse. Beds committed only through the logged decision.	Audit capacity (1 or unlimited reviews/turn)
	P4: No private data leaks between agents. A downstream agent can't be handed demographic detail it doesn't need (not about leaks to end users, which is a different check).	
	P5: Full traceability. The decision must be fully reconstructable from the record.	

To actually run this at scale, we didn't click through the app by hand: a script calls the same API the app itself uses, once per agent per case, and logs every decision straight to a database, then repeats that across every combination of caseload curve, bed count, and audit capacity we wanted to test. That produced 192 full episodes and 2,304 resolved case pairs, 0 failed runs, all on GPT-4o-mini rather than a larger model, mainly to keep the API cost of running that many episodes manageable.

2. Results

30.0% of unfair outcomes leave no trace anywhere in the decision chain: never flagged by an allocator, never caught by an audit. That share depends heavily on audit load: 18.4% when the auditor isn't overloaded, 43.8% when it is.

Table 1 shows that splitting the decision across roles barely moves the bias rate itself:


TABLE 1. BIAS RATE BY CONDITION.

CONDITION	BIAS RATE	N
Control	6.9%	79 / 1,152
Multi-agent	6.1%	70 / 1,152
Difference		+0.8pp, 95% CI [-1.2, +2.8], p = 0.498 (n.s.)

Splitting the decision across roles doesn't change how often identical twins get different outcomes. Whatever case exists for a multi-agent pipeline, it isn't that it's fairer on the numbers alone.

Figure 1 separates two different failure modes: whether an overloaded auditor gets worse at catching bias, or simply looks at fewer cases in the first place. Coverage is whether a case gets looked at at all; judgment is whether the bias gets caught once it is.

FIGURE 1. COVERAGE VS. JUDGMENT, FULL AUDIT CAPACITY VS. AN OVERLOADED AUDITOR.



Coverage, whether a case gets reviewed at all, collapses from 100.0% to 65.6% under a 1-review-per-turn cap ($p < 0.001$). Judgment quality on the cases that *are* reviewed barely moves: 81.6% vs 85.7% ($p = 1.000$, not significant, direction even reverses). The auditor doesn't get worse under pressure, it just sees less.

If the real bottleneck is which cases get looked at, the fix should target that directly. So we tried one concrete change: instead of auditing decisions in the order they arrive (a FIFO queue, first in, first out), audit the riskiest ones first (a risk-priority queue), meaning any decision whose rationale names the patient's nationality, religion, or other protected trait, or that ended in a denial, jumps to the front. Table 2 shows what that reordering alone recovers:

TABLE 2. FIFO VS. RISK-PRIORITY AUDIT QUEUE, SAME CASES AND AUDIT CAPACITY, PAIRED BY SEED.

METRIC	FIFO QUEUE	RISK-PRIORITY QUEUE	P
Coverage	65.6%	91.7%	0.028
Catch rate	56.2%	79.2%	0.092
Silent-bias rate	43.8%	20.8%	0.092

Silent-bias rate is the same "leaves no trace anywhere" measure from the start of this section (30.0% overall), broken out here by queue strategy instead of by audit capacity. Coverage recovers significantly. The silent-bias and catch-rate gains in Table 2 point the same direction but aren't confirmed at this sample size, and silent-bias here is the mathematical mirror of catch rate, not an independent check.

3. Discussion

The clearest result here isn't about fairness. Splitting the decision across roles doesn't change how often the model treats identical twins differently (Table 1). That's a narrower claim than it might sound: KillBench [1] itself found that process choices can matter a lot, switching the output format from structured to free text shifted both the refusal rate and the bias it measured, so this isn't a case of citing them to say bias is fixed regardless of process. What we varied here (splitting one decision across roles, not the output format) is a different kind of process change, and on that specific axis, the rate doesn't move. What does change is detection: across all five policies, not just the bias one, a single agent auditing its own work essentially never contradicts itself (unsurprising once you notice it's the same agent, in the same call, so there's no real separation between deciding and reviewing), while a separate auditor catches real problems at rates from 7% to 62%, depending on the policy. That is the same shape of failure PBSuite [2] documents under adversarial pressure, except here nothing adversarial is happening. The pressure is purely operational: caseload, bed scarcity, audit capacity.

That detection advantage erodes under load (Figure 1), but the erosion is a scheduling problem, not a judgment one: coverage collapses from 100.0% to 65.6% while the quality of review on cases actually seen barely moves. That is a different failure than the one DeZoort and Braun [5] describe in human auditors, where time pressure degrades judgment directly. Here the auditor doesn't get worse at its job, it just does the job on fewer cases, closer to a queueing failure than a cognitive one. This also lines up with MAST's [4] task-verification failure category: the failure here isn't in any single agent's reasoning, it's in whether the check on that reasoning actually happens. It's not one bad agent either: silent-bias cases are spread evenly across the three allocator instances (chi-square = 0.99, $p = 0.61$) and split exactly 50/50 across the two auditors, so there's no single culprit to swap out. That is probably not specific to this setup. Any team that adds independent oversight to an agent pipeline will hit some version of the same trade-off, where the auditor is only as good as its queue, and a silent coverage gap under load can look, from the outside, like the system is working when it isn't.

The bed-scarcity mechanic is adapted from GovSim's [3] shared-resource setup. GovSim's agents typically fail by over-extracting a shared resource out of short-term self-interest; ours don't fail that way; bed allocations track severity about as well under pressure as without it (Table 1). What runs out instead is a different resource entirely, audit attention, and that shortage doesn't show up anywhere in the allocation numbers, only in how much of the record ever gets checked (Figure 1).

This study also cut corners for practical reasons. Running one model across 192 episodes plus the follow-up experiment already added up in API cost, so we didn't compare multiple models against each other, and we didn't push the episode count past what's reported here. Both would be the obvious next step before trusting these numbers beyond GPT-4o-mini specifically.

4. Limitations

This is one model, GPT-4o-mini. KillBench itself found bias varies by model, so nothing here generalizes past it. Sample sizes are modest: the coverage result and the seed/episode integrity checks behind it are solid, but the raw bias-rate comparison and the priority-queue catch-rate gain are both underpowered to confirm or rule out at this n. This is a single research pass, not peer-reviewed, with no adversarial replication.

5. References

1. **KillBench** (White Circle): the single-agent demographic-bias finding this study extends to a multi-agent pipeline.
2. **PBSuite** (Varshney et al., NVIDIA): policy compliance under adversarial pressure; this study substitutes operational pressure (caseload, capacity) instead.
3. **GovSim / "Cooperate or Collapse"** (Piatti et al., ETH Zürich): source of the shared, depletable resource (beds) mechanic used here.
4. **MAST** (Cemri et al.): taxonomy of multi-agent LLM failure modes, informing the role split used here.
5. **DeZoort (1998) & Braun (2000)**: audit time-pressure literature; motivated the judgment-degradation hypothesis that this data does not confirm.